
File System Management

COP-5614 Operating System
Group 9:

Md Shahadat Iqbal Md Abdullah Al Mamun Vitalii Stebliankin



Contributions

| Name | Contribution |
|-----------------------------|---|
| Md Abdullah Al Mamun | Bitmap and File name checking <ul style="list-style-type: none">• <code>bitmap_init()</code>• <code>bitmap_first_unused()</code>• <code>bitmap_reset()</code>• <code>illegal_file_name()</code> |

Contributions

| Name | Contribution |
|-----------------------------|---|
| Md Abdullah Al Mamun | Bitmap and File name checking <ul style="list-style-type: none">• <code>bitmap_init()</code>• <code>bitmap_first_unused()</code>• <code>bitmap_reset()</code>• <code>illegal_file_name()</code> |
| Vitalii Stebliankin | Remove <ul style="list-style-type: none">• <code>File Unlink()</code>• <code>Dir Unlink()</code>• <code>remove_file_or_directory()</code>• <code>remove_inode()</code> |

Contributions

| Name | Contribution |
|-----------------------------|---|
| Md Abdullah Al Mamun | Bitmap and File name checking <ul style="list-style-type: none">• <code>bitmap_init()</code>• <code>bitmap_first_unused()</code>• <code>bitmap_reset()</code>• <code>illegal_file_name()</code> |
| Vitalii Stebliankin | Remove <ul style="list-style-type: none">• <code>File Unlink()</code>• <code>Dir Unlink()</code>• <code>remove_file_or_directory()</code>• <code>remove_inode()</code> |
| Md Shahadat Iqbal | Read/Write <ul style="list-style-type: none">• <code>File_Read()</code>• <code>File_Write()</code>• <code>File_Seek()</code>• <code>Dir_Read()</code>• <code>Dir_Size()</code> |

Bitmap Initialization

- **Set all '0'**
 - `memset()`
- **Set first *nbits* '1'**

Buffer = 0

1 (OR)

1 (*nbits*)

- **Features**
 - If *nbits* is zero
 - No further checking

```
static void bitmap_init(int start, int num, int nbits){
    int bit, byte, sector;

    //initializing 0 in bitmap of every sector
    int flag1 = 1;
    sector = 0;
    while(sector < num){
        char buffer_bitmap[SECTOR_SIZE];          // bitma
        memset(buffer_bitmap, 0, SECTOR_SIZE);    // fill
        for (byte = 0; byte < SECTOR_SIZE; byte++){ // loop
            if(!flag1) {
                break;
            }
            else {
                for (bit = 7; bit >= 0; bit--){          // loo
                    if (nbits-- > 0){                    // chekc
                        buffer_bitmap[byte] |= (1 << bit);
                    }
                }
                else{
                    flag1 = 0;
                    break;
                } // set 1 for first nbits (bitwis
            }
        }
        //update on the disk
        Disk_Write(start+sector, buffer_bitmap);
        sector++;
    }
}
```

Find First Unused

- **Bitwise AND**

Buffer= 0

1 (AND)

0

- **Flip “0” to “1”, return loc**

Buffer = 0

1 (OR)

1

- **Features**

- If $id > nbits$
- No further checking

```
static int bitmap_first_unused(int start, int num, int nbits){
    int id = 0; //starting index is zero by default
    char check_bit_empty;
    char buffer_bitmap[SECTOR_SIZE];
    int bit, byte, sector;

    // checking every sectors to find the first unused memory
    sector = 0;
    while(sector < num){
        // looping each sector
        Disk_Read(start+sector, buffer_bitmap); // pull current sector
        for (byte = 0; byte < SECTOR_SIZE; byte++){ // looping each byte
            for (bit = 7; bit >= 0; bit--){ // looping each bit
                check_bit_empty = (buffer_bitmap[byte] >> bit) & 1; // check if bit is empty

                if (check_bit_empty == 0){ // if used bit
                    buffer_bitmap[byte] |= 1 << bit;
                    Disk_Write(start+sector, buffer_bitmap); // update the buffer
                    return id; // return the index of the first unused bit
                }
            }
            id++;
        }
        if (id > nbits){ // return -1 if the index exceeds the number of bits
            return -1;
        }
        // increment index by 1
        sector++;
    }
    return -1; // return -1 if no unused bit found
}
```

Bitmap Reset

- Pulling current bitmap
- Set ibit to "0"
 - Bitwise AND

```
static int bitmap_reset(int start, int num, int ibit) {  
    char buffer_bitmap[SECTOR_SIZE];  
    int bit, byte, sector;  
  
    // loop through each sector in bitmap  
    sector = 0;  
    while(sector < num){  
        Disk_Read(start+sector, buffer_bitmap);  
        for (byte = 0; byte < SECTOR_SIZE; byte++){ // che  
            for(bit = 7; bit >= 0; bit--){ // che  
                if(!ibit){ // onc  
                    buffer_bitmap[byte] &= ~(1 << bit); // Let  
                    Disk_Write(start+sector, buffer_bitmap);  
                    return 0;  
                }  
                ibit--;  
            }  
        }  
        sector++;  
    }  
    return -1;  
}
```

Illegal filename

- **Valid characters**

- letters
- dots, dashes, and underscore
- numbers

- **Length**

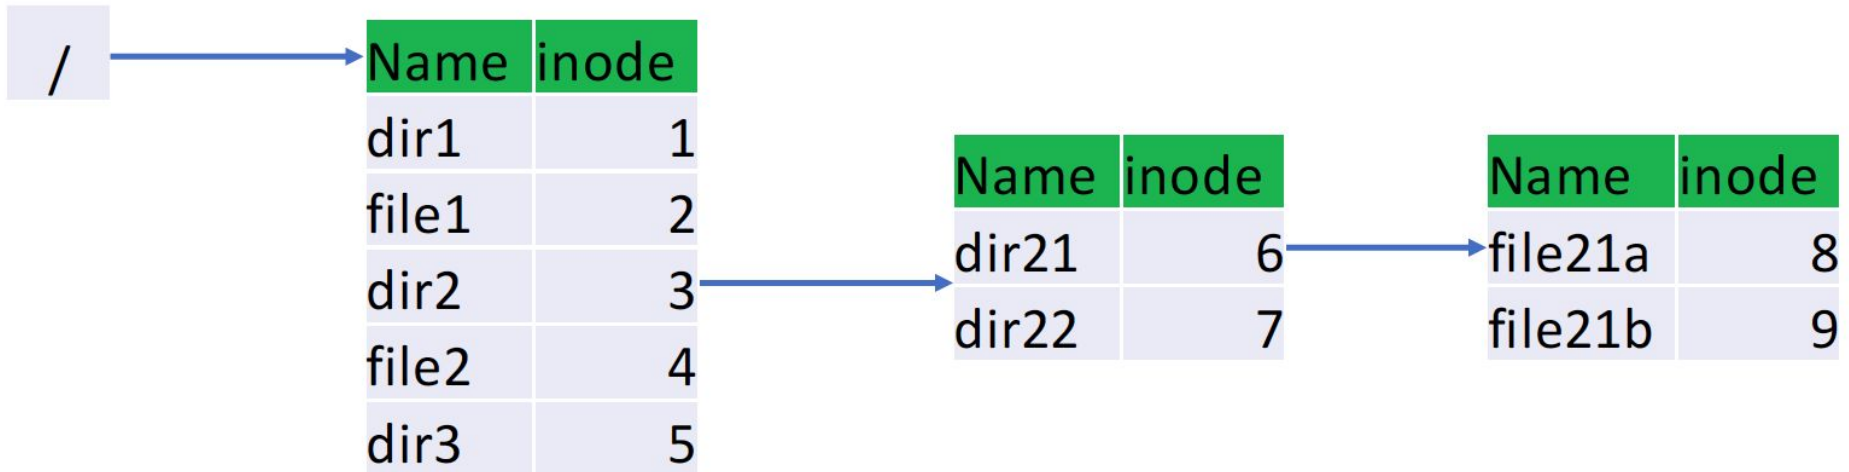
- MAX_NAME-1

```
static int illegal_filename(char* name) {  
    int i;  
    char end[] = "-_.";   
  
    //checking character by character  
    for(i = 0; i < strlen(name); i++){  
        if(!isdigit(name[i]) && !isalpha(name[i]) && !strchr(end, name[i])){  
            printf("Illegal character found! %c\n", name[i]);  
            return 1;  
        }  
    }  
  
    //checking length  
    if(strlen(name) >= MAX_NAME-1){  
        printf("Bad Length\n");  
        return 1;  
    }  
  
    //return false otherwise  
    return 0;  
}
```


Demo: Bitmap

1. Simple test
 - a. Create file/directory
 - b. Remove or unlink file/directory
2. Customized test: **5 directories and 4 files**

A. Initial inode structure




Remove File or Directory

- Remove iNode from the disk and update parent iNode



- Remove the content of file from the disk



- Update bitmaps

iNode bitmap

| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|---|--------------|---|---|---|---|---|---|---|----|----|----|----|----|
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 0 | 0 | 0 | 0 | 0 |

0

Disk Sector bitmap

| 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|---|---|---|---|---|--------------|---|---|---|----|----|----|----|----|
| 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 1 | 0 | 1 | 1 | 0 |

0

remove_inode() - Part 1

- Load iNode from the disk
- Check for errors
- Remove content from the disk

```
int remove_inode(int type, int parent_inode, int child_inode)
{
    // This function is a modification to add_inode() that removes child inode
    dprintf("... removing inode %d\n", child_inode);
    //-----
    // Load Child iNode from the disk
    //-----
    // load the disk sector containing the child inode
    int inode_sector = INODE_TABLE_START_SECTOR+child_inode/INODES_PER_SECTOR;
    char inode_buffer[SECTOR_SIZE];
    if(Disk_Read(inode_sector, inode_buffer) < 0){
        dprintf("Error: can't read inode from disk sector");
        return -1;
    }
    //dprintf("%s\n", inode_buffer);
    dprintf("... load inode table for child inode from disk sector %d\n", inode_sector);

    // get the child inode
    int inode_start_entry = (inode_sector-INODE_TABLE_START_SECTOR)*INODES_PER_SECTOR;
    int offset = child_inode-inode_start_entry;
    assert(0 <= offset && offset < INODES_PER_SECTOR);
    inode_t* child = (inode_t*)(inode_buffer+offset*sizeof(inode_t));

    //-----
    // Check for errors
    //-----
    // Check if the type is right
    if(child->type != type){
        dprintf("Error: the actual type of inode doesn't match with parameter specified in function remove_inode()");
        return -3;
    }

    //Check if directory is not empty
    if(child->type==1 && child->size!=0){
        dprintf("Error: directory is not empty");
        osErrno = E_DIR_NOT_EMPTY;
        return -2;
    }

    //-----
    // Remove the file content from the disk
    //-----
    char buf_sector[SECTOR_SIZE];
    int sector_address;
    if(child->size>0){
        dprintf("%s\n", "... Removing file data from the disk");
        for(int i; i<MAX_SECTORS_PER_FILE; i++){
            sector_address = child->data[i];
            if(sector_address>0){
                //remove from sector
                if(Disk_Read(sector_address, buf_sector)<0) return -1; //try to read the sector
                memset(buf_sector, 0, SECTOR_SIZE); //set buffer to 0
                Disk_Write(sector_address, buf_sector); //write 0s to the sector
                bitmap_reset(SECTOR_BITMAP_START_SECTOR, SECTOR_BITMAP_SECTORS, sector_address); //reset bitmap for the current sector
            }
        }
    }
    else{
        dprintf("%s\n", "... The directory/file is empty");
    }
}
```


remove_inode() - Part 2

- Remove iNode from the disk
- Reset iNode bitmap
- Read parent iNode to the disk

```
301 //
302 // Remove iNode from the disk sector
303 //
304 dprintf("%s\n", "... Removing inode from the disk");
305 memset(child, 0, sizeof(inode_t)); //Delete child inodes from memory
306 // memset(inode_buffer, 0, SECTOR_SIZE);
307 if(Disk_Write(inode_sector, inode_buffer)<0) return -1;
308 //
309 // Reset inode bitmap
310 //
311 bitmap_reset(INODE_BITMAP_START_SECTOR, INODE_BITMAP_SECTORS, child_inode); //reset i-node bitmap
312 // Remove link from parent inode to child inode
313 //
314 // Read the parent iNode from the disk
315 //
316 // get the disk sector containing the parent inode
317 inode_sector = INODE_TABLE_START_SECTOR+parent_inode/INODES_PER_SECTOR;
318 if(Disk_Read(inode_sector, inode_buffer) < 0) return -1;
319 dprintf("... load inode table for parent inode %d from disk sector %d\n",
320 parent_inode, inode_sector);
321 //
322 // get the parent inode
323 inode_start_entry = (inode_sector-INODE_TABLE_START_SECTOR)*INODES_PER_SECTOR;
324 offset = parent_inode-inode_start_entry;
325 assert(0 <= offset && offset < INODES_PER_SECTOR);
326 inode_t* parent = (inode_t*)(inode_buffer+offset*sizeof(inode_t));
327 dprintf("... get parent inode %d (size=%d, type=%d)\n",
328 parent_inode, parent->size, parent->type);
329 //
330 // Check if parent iNode is a directory
331 //
332 if(parent->type != 1){
333 dprintf("... error: parent inode is not directory\n");
334 return -2;
335 }
```

remove_inode() - Part 3

- Update parent inode



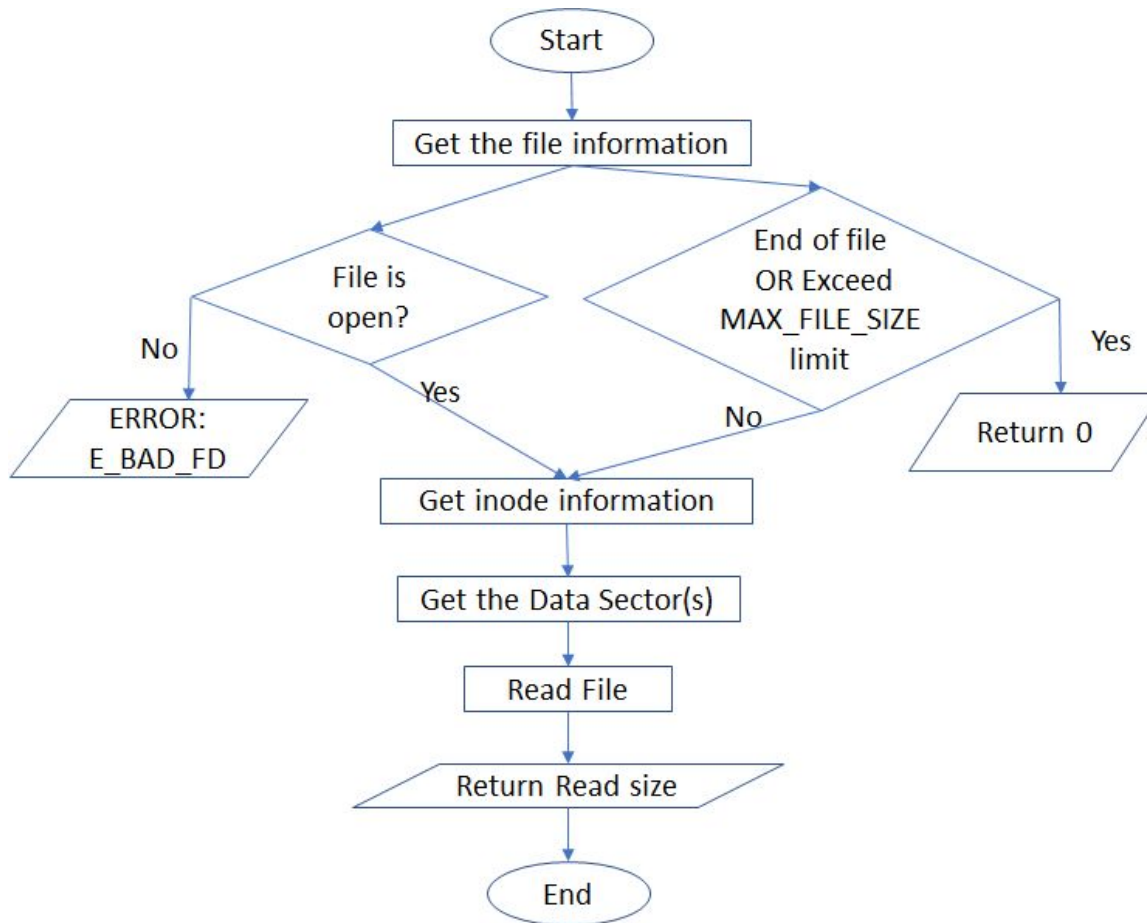
```
336 //
337 // Update parent inode and dirent table
338 //
339 int nentries = parent->size; // remaining number of directory entries
340 int idx = 0;
341 int found_flag=0;
342 while(nentries > 0) {
343     char dirent_buffer[SECTOR_SIZE]; // cached content of directory entries
344     if(Disk_Read(parent->data[idx], dirent_buffer) < 0) return -2;
345     for(int i=0; i<DIRENTS_PER_SECTOR; i++) {
346         if(i>nentries) break;
347         dirent_t* tmp_dirent = (dirent_t*)(dirent_buffer + i*sizeof(dirent_t));
348
349         if(tmp_dirent->inode==child_inode){//dirent found
350             dprintf("... Corresponding dirent found\n");
351             char dirent_buffer2[SECTOR_SIZE];
352
353             int group = parent->size/DIRENTS_PER_SECTOR;
354             if(Disk_Read(parent->data[group], dirent_buffer2) < 0) return -2;
355             int start_entry = group*DIRENTS_PER_SECTOR;
356             offset = parent->size-start_entry;
357             dirent_t* tmp_dirent1 = (dirent_t*)(dirent_buffer2+offset*sizeof(dirent_t));
358
359             strncpy(tmp_dirent->fname, tmp_dirent1->fname, MAX_NAME);
360             tmp_dirent->inode = tmp_dirent1->inode;
361             memset(tmp_dirent1, 0, sizeof(dirent_t));
362             parent->size--;
363             if(Disk_Write(inode_sector, inode_buffer)<0) return -1;//update parent inode
364             if(Disk_Write(parent->data[group], dirent_buffer2)) return -1;//update dirent
365             if(Disk_Write(parent->data[idx], dirent_buffer)) return -1;//update dirent
366             found_flag=1;
367         }
368     }
369     idx++; nentries -= DIRENTS_PER_SECTOR;
370 }
371 if(found_flag==1) return 0;
372 else{
373     dprintf("... could not find dirent\n");
374     return -1; // not found
375 }
376 }
```

Read/Write

1. File_Read()
2. File_Write()
3. File_Seek()
4. Dir_Read()
5. Dir_Size()

File Read

```
int File_Read(int fd, void* buffer, int size)
```



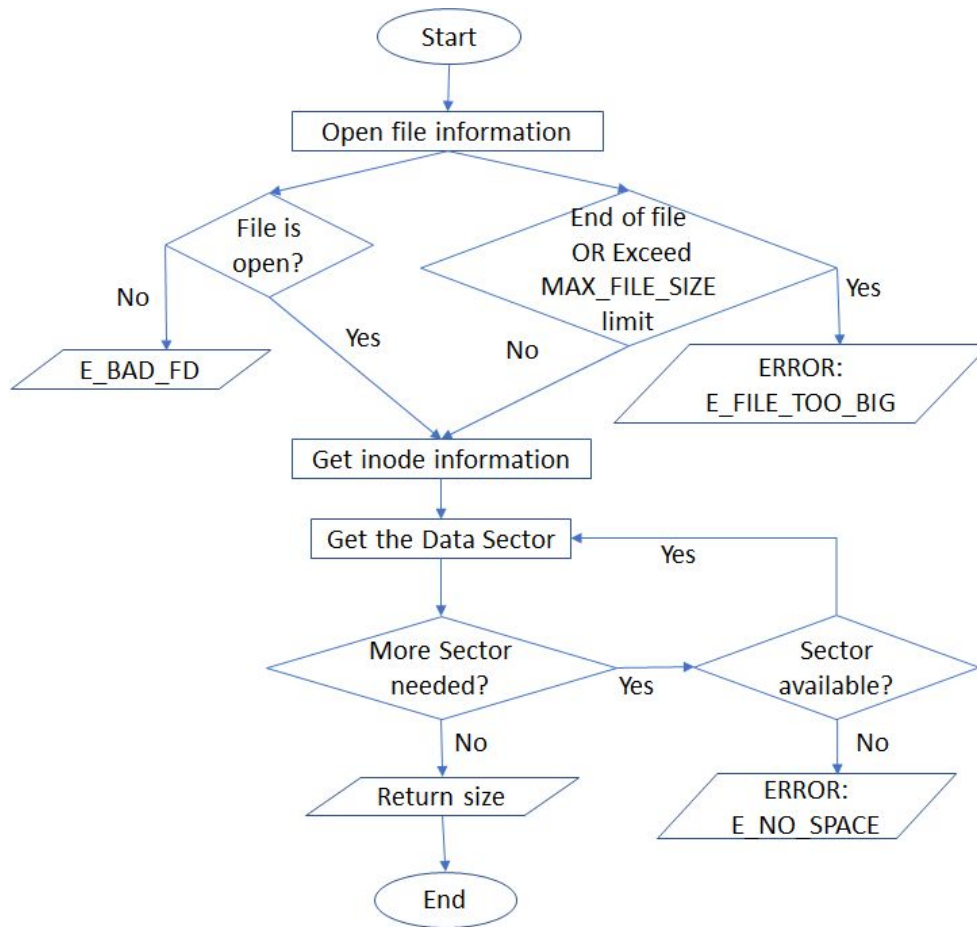
```

int File_Read(int fd, void* buffer, int size)
{
    int reading_part = 0; // file is read and kept a copy into reading_part
    open_file_t openFileInfo = open_files[fd]; // Information about the file fd
    // first check the file is open or not
    if (openFileInfo.inode == 0)
    {
        dprintf("...file not open");
        oserrno = E_BAD_FD;
        return -1;
    }
    // load the disk sector containing the inode
    int inode_sector = INODE_TABLE_START_SECTOR + openFileInfo.inode / INODES_PER_SECTOR;
    char inode_buffer[SECTOR_SIZE];
    if (Disk_Read(inode_sector, inode_buffer) < 0) {
        oserrno = E_GENERAL;
        return -1;
    }
    dprintf("... load inode table for inode from disk sector %d\n", inode_sector);
    // get the inode
    int inode_start_entry = (inode_sector - INODE_TABLE_START_SECTOR) * INODES_PER_SECTOR;
    int offset = openFileInfo.inode - inode_start_entry;
    assert(0 < offset && offset < INODES_PER_SECTOR);
    inode_t *child_inode = (inode_t *) (inode_buffer + offset * sizeof(inode_t));
    int sector_read = openFileInfo.pos / SECTOR_SIZE; // The sector which need to read
    // Check if the position is at the end of file, or data at that sector does not exists
    if (openFileInfo.pos == MAX_FILE_SIZE || !child_inode->data[sector_read]) {
        return 0;
    }
    char sectorBuffer[SECTOR_SIZE];
    // check if there any problem to read the data
    if (Disk_Read(child_inode->data[sector_read], sectorBuffer) < 0) {
        return -1;
    }
    dprintf("... load disk sector %d\n", child_inode->data[sector_read]);
    void *reading_pos = sectorBuffer + (open_files[fd].pos - (openFileInfo.pos / SECTOR_SIZE) * SECTOR_SIZE);
    // Read the data from the file until maximum file is already read completely or
    // requested read size is already read or all the data has been accessed
    while (open_files[fd].pos < MAX_FILE_SIZE && size > 0 && child_inode->data[sector_read]) {
        if (size >= SECTOR_SIZE) {
            dprintf("... Reading BYTE %d\n", open_files[fd].pos);
            memcpy(buffer, reading_pos, SECTOR_SIZE);
            size -= SECTOR_SIZE;
            buffer += SECTOR_SIZE;
            reading_part += SECTOR_SIZE;
            open_files[fd].pos += SECTOR_SIZE;
            sector_read++;
            if (open_files[fd].pos < MAX_FILE_SIZE) { // check if file read reach to maximum file size
                if (Disk_Read(child_inode->data[sector_read], sectorBuffer) < 0) {
                    return -1;
                }
                reading_pos = sectorBuffer;
            } else {
                return reading_part;
            }
        } else {
            dprintf("... Reading BYTE %d\n", open_files[fd].pos);
            memcpy(buffer, reading_pos, (size_t) size);
            reading_part += size;
            open_files[fd].pos += size;
            return reading_part;
        }
    }
    return reading_part;
}

```


File Write

```
int File_Write(int fd, void* buffer, int size):
```



```

/* YOUR CODE */
int File_Write(int fd, void* buffer, int size)
{
    int remaining_writing = size;
    // Have the information about the file (fd) where we have to write the buffer
    open_file_t openFileInfo = open_files[fd];
    // First check the file is open or not
    if (openFileInfo.inode == 0) {
        dprintf("... file not open");
        osErrno = E_BAD_FD; // If not open set the osError to E_BAD_FD
        return -1;
    }

    // check if the file have enough space to write
    if ((openFileInfo.size + size) > MAX_FILE_SIZE) {
        dprintf("... file is too big.");
        osErrno = E_FILE_TOO_BIG;
        return -1;
    }

    // load the disk sector containing the inode
    int inode_sector = INODE_TABLE_START_SECTOR + openFileInfo.inode / INODES_PER_SECTOR;
    char inode_buffer[SECTOR_SIZE];
    if (Disk_Read(inode_sector, inode_buffer) < 0) {
        osErrno = E_GENERAL;
        return -1;
    }
    dprintf("... load inode table for inode from disk sector %d\n", inode_sector);

    // get the inode
    int inode_start_entry = (inode_sector - INODE_TABLE_START_SECTOR) * INODES_PER_SECTOR;
    int offset = openFileInfo.inode - inode_start_entry;
    assert(0 <= offset && offset < INODES_PER_SECTOR);
    inode_t *fileInode = (inode_t *) (inode_buffer + offset * sizeof(inode_t));

    int sectors_need = (size / SECTOR_SIZE) + 1; // number of sectors needed to write the buffer

    // start the process of writing
    for (int i = 0; i < sectors_need; i++) {
        char sector[SECTOR_SIZE];
        int sector_ID = bitmap_first_unused(SECTOR_BITMAP_START_SECTOR, SECTOR_BITMAP_SECTORS, SECTOR_BITMAP_SIZE);

        // Check if there is any sector left for writing
        if (sector_ID < 0) {
            osErrno = E_NO_SPACE;
            dprintf("... no space left.");
            Disk_Write(INODE_TABLE_START_SECTOR + (openFileInfo.inode / INODES_PER_SECTOR), inode_buffer);
            return -1;
        }
        fileInode->data[i] = sector_ID;
        Disk_Read(sector_ID, sector);

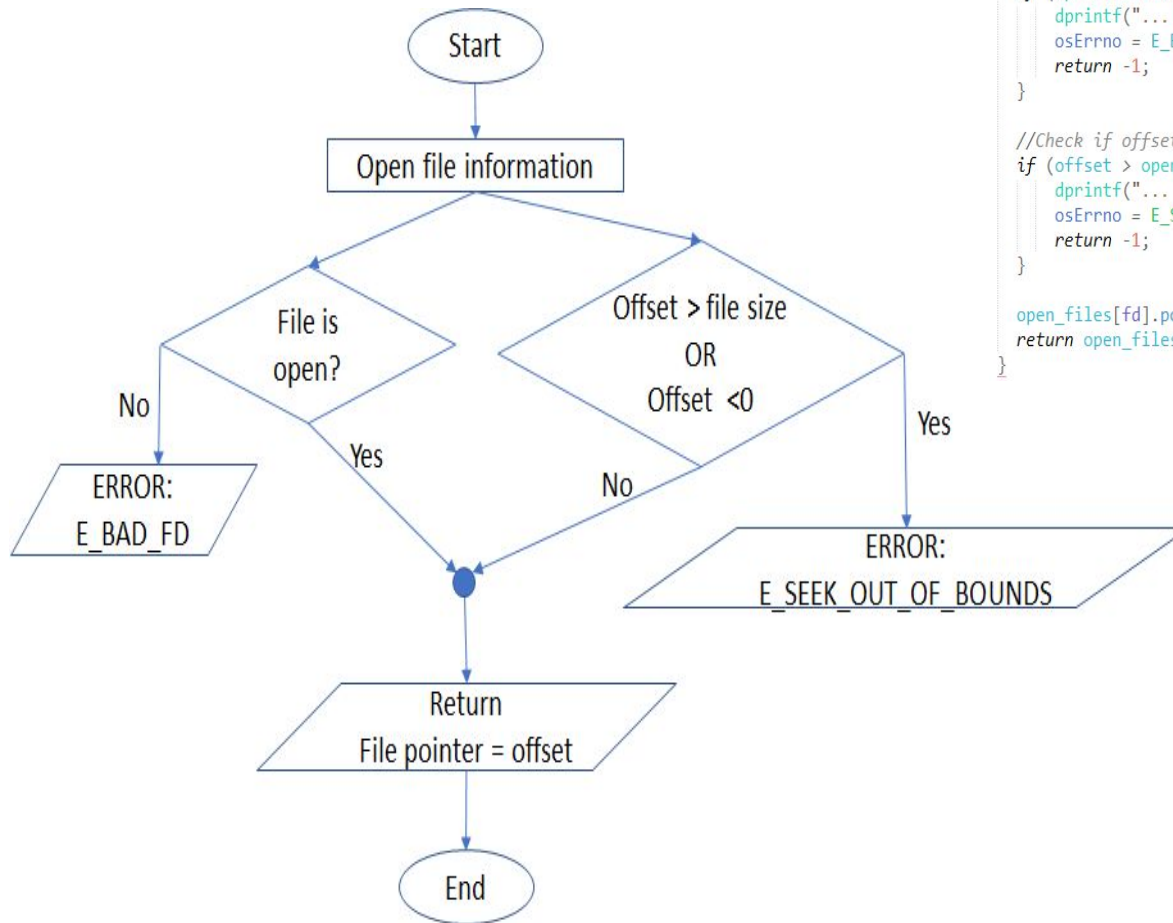
        // write buffer to sector. Buffer might be bigger than sector
        // need to request more sectors
        if (remaining_writing > SECTOR_SIZE) {
            remaining_writing -= SECTOR_SIZE;
            memcpy(sector, buffer, SECTOR_SIZE);
            buffer += SECTOR_SIZE;
            fileInode->size += SECTOR_SIZE;
        } else {
            memcpy(sector, buffer, (size_t) remaining_writing);
            buffer += remaining_writing;
            fileInode->size += remaining_writing;
            remaining_writing = 0;
        }
        openFileInfo.size = fileInode->size;
        openFileInfo.pos = fileInode->size;
        Disk_Write(sector_ID, sector);
    }
    Disk_Write(INODE_TABLE_START_SECTOR + (openFileInfo.inode / INODES_PER_SECTOR), inode_buffer);
    return size - remaining_writing;
}

```


File Seek

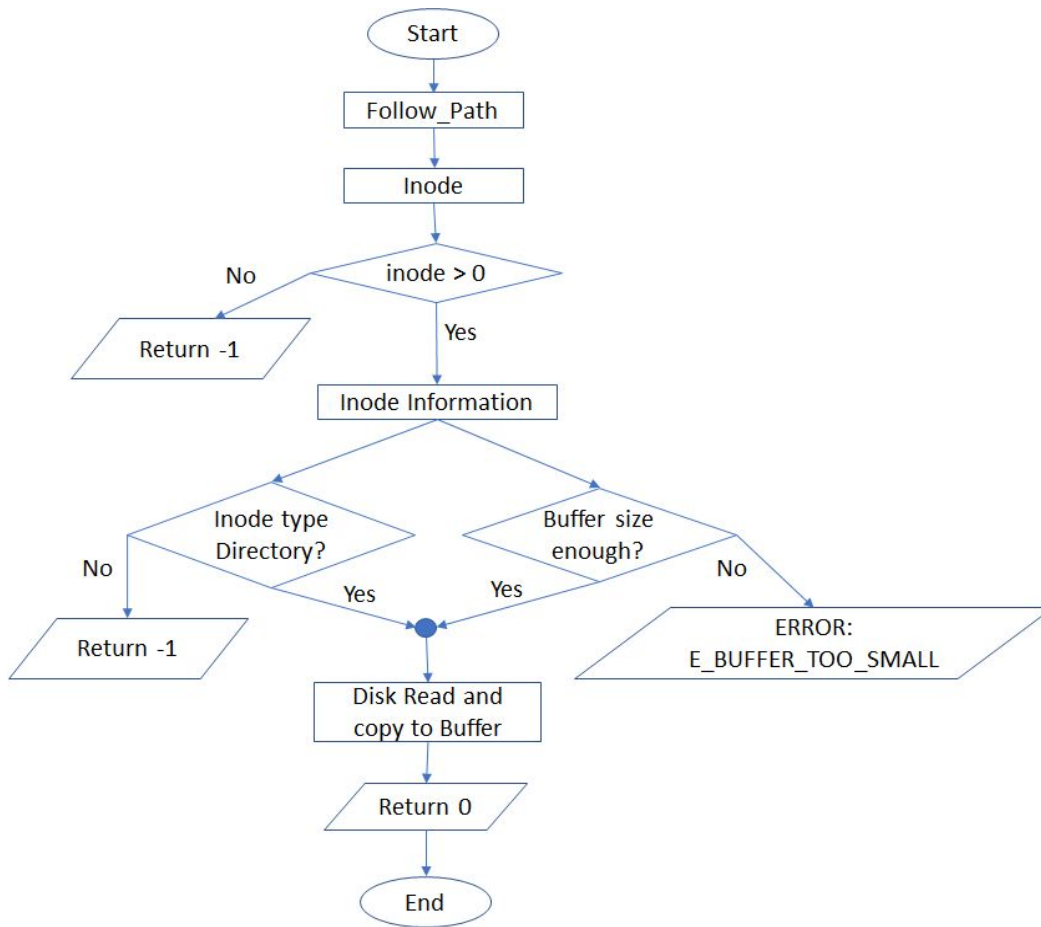
```
int File_Seek(int fd, int offset):
```

```
int File_Seek(int fd, int offset){  
    open_file_t openFileInfo = open_files[fd];  
  
    if (openFileInfo.inode == 0) { // check if the file is open or not  
        dprintf("... file not open");  
        osErrno = E_BAD_FD;  
        return -1;  
    }  
  
    //Check if offset is larger than the size of the file or offset is negative  
    if (offset > openFileInfo.size || offset < 0) {  
        dprintf("... offset is out of bound");  
        osErrno = E_SEEK_OUT_OF_BOUNDS;  
        return -1;  
    }  
  
    open_files[fd].pos = offset;  
    return open_files[fd].pos;  
}
```



Directory Read

```
int Dir_Read(char* path, void* buffer, int size)
```



```

int Dir_Read(char* path, void* buffer, int size) {
    int parent_inode; //inode of the directory that need to read
    follow_path(path, &parent_inode, NULL);

    if (parent_inode >= 0) { //check the directory is found or not
        // Load the disk sector containing the inode
        int inode_sector = INODE_TABLE_START_SECTOR+parent_inode/INODES_PER_SECTOR;
        char inode_buffer[SECTOR_SIZE];
        if(Disk_Read(inode_sector, inode_buffer) < 0) {
            osErrno = E_GENERAL;
            return -1;
        }
        dprintf("... Load inode table for inode from disk sector %d\n", inode_sector);

        // get the inode
        int inode_start_entry = (inode_sector-INODE_TABLE_START_SECTOR)*INODES_PER_SECTOR;
        int offset = parent_inode-inode_start_entry;
        assert(0 <= offset && offset < INODES_PER_SECTOR);
        inode_t* directory_inode = (inode_t*)(inode_buffer+offset*sizeof(inode_t));
        dprintf("... inode %d (size:%d, type:%d)\n", parent_inode, directory_inode->size, directory_inode->type);

        // check whether the inode type is directory
        if(directory_inode->type != 1) {
            dprintf("... error: '%s' is not a Directory\n", path);
            osErrno = E_GENERAL;
            return -1;
        }

        //check if the read buffer is larger enough to hold the details
        if(directory_inode->size*sizeof(dirent_t) > size){
            dprintf("ERROR: type-%d inode-%d size-%d givesize-%d\n", directory_inode->type, parent_inode, directory_inode->size, size);
            osErrno = E_BUFFER_TOO_SMALL;
            return -1;
        }


        int sector, i;
        int increase_size = 0; // increase the buffer size by this amount
        char dirent_buffer[SECTOR_SIZE];

        for(sector = 0; sector < MAX_SECTORS_PER_FILE; sector++){ //iterate over each sector
            if(directory_inode->data[sector]){
                Disk_Read(directory_inode->data[sector], dirent_buffer);
                for(i = 0; i < DIRENTS_PER_SECTOR; i++){ // for each directory read it and copy it to the buffer
                    dirent_t* dirent = (dirent_t*)(dirent_buffer+i*sizeof(dirent_t));
                    if(dirent->inode){
                        memcpy(buffer+increase_size, (void*)dirent, sizeof(dirent_t));
                        increase_size += sizeof(dirent_t);
                    }
                }
            }
        }

        dprintf("%d\n", directory_inode->size);
        return directory_inode->size;
    } else {
        dprintf("... directory '%s' is not found\n", path);
        return 0;
    }
}
  
```

Directory Size

```
int Dir_Size(char* path)
```



```
int Dir_Size(char* path){
    int parent_inode; //inode of the directory that need to read
    follow_path(path, &parent_inode, NULL);

    if (parent_inode >= 0) { //check the directory is found or not

        // Load the disk sector containing the inode
        int inode_sector = INODE_TABLE_START_SECTOR+parent_inode/INODES_PER_SECTOR;
        char inode_buffer[SECTOR_SIZE];
        if(Disk_Read(inode_sector, inode_buffer) < 0) {
            osErrno = E_GENERAL;
            return -1;
        }
        dprintf("... load inode table for inode from disk sector %d\n", inode_sector);

        // get the inode
        int inode_start_entry = (inode_sector - INODE_TABLE_START_SECTOR) * INODES_PER_SECTOR;
        int offset = parent_inode - inode_start_entry;
        assert(0 <= offset && offset < INODES_PER_SECTOR);
        inode_t *directory_inode = (inode_t *) (inode_buffer + offset * sizeof(inode_t));
        dprintf("... inode %d (size=%d, type=%d)\n", parent_inode, directory_inode->size, directory_inode->type);

        // check weather the inode type is directory
        if (directory_inode->type != 1) {
            dprintf("... error: '%s' is not a directory\n", path);
            osErrno = E_GENERAL;
            return -1;
        }
        dprintf("... RETURNING SIZE: '%d' \n", (int) (directory_inode->size * sizeof(dirent_t)));

        return (int) (directory_inode->size * sizeof(dirent_t));
    } else {
        dprintf("... directory '%s' is not found\n", path);
        return 0;
    }
}
```