

Project Report: File System Management

COP-5614 Operating System

Group:

Md Shahadat Iqbal - 4887079

Md Abdullah Al Mamun - 6144422

Vitalii Stebliankin - 6102408

Professor: Liting Hu

Submission date: December 2, 2019

1 Introduction

The project focus on developing a file management system. The main objectives of this project are:

- Gain a details understanding on how does file system work, specifically the directory hierarchy and storage management.
- Gain an understanding on some of the performance issues file systems must deal with.

In this project, we have built the user-level library, *libFS*, that can simulate some of the functions of a file management system.

1.1 Contribution

We divided the implementation into three parts and each group member is assigned to each part as follows.

- Bitmap and miscellaneous by Md Abdullah Al Mamun
- Read/Write of file or directory by Md Shahadat Iqbal
- Remove file/directory/inode by Vitalii Stebliankin

Assigned member has full contribution for designing and developing related functions. Code developing, checking possible test cases, writing a report, and preparing the presentation were performed under the active collaboration and productive discussions. Even though each of us was responsible for its part, we helped each other with the architecture and code design.

2 Implementation

2.1 Bitmap

Memory space in the disk is limited. Thus, managing free space is crucial part of any file system. File management includes keeping track of unused disk blocks, allocating space for

newly created files and reuse the released space when a file is removed. For handling such operation, Bitmap or bit vector is used that contains the index of memory space. The disk is divided into units for bitmap. These units are range from a few bytes to several kilobytes. Each memory unit is associated with a bit in the bitmap. The value of a bit could be either 0 or 1 where 0 indicates the unit is free and 1 indicates that it is occupied. In this file system, two bitmaps has been created for inode and files. In this section, implementation design of bitmap initialization, finding the first unused bit and reset a bit has been described.

2.1.1 Initialization

bitmap_init() function has three parameters, *start*, *num*, and *nbits*. This function is initializing zero value in the bit vector with *num* of sectors starting from *start* sector except that the first *nbits* number of bits. The number of first *nbits* are set to one. The code snippet of this function is shown in Fig 1. A bitmap buffer of character type is created because each character size is 8 bit which is equal to 1 byte. First each bit of the every sector are set to 0 then we changed first *nbits* from 0 to 1 in line# 16-18. A single 1 is shifted to the left shifted and then a bitwise OR operation is performed to make them 1 for number of first *nbits*. *flag1* is used to avoid unnecessary loop execution means if *nbits* are already set to 1, then no need to check remaining sector/bytes. This improves the efficiency.

```

1  static void bitmap_init(int start, int num, int nbits){
2      int bit, byte, sector;
3
4      //initializing 0 in bitmap of every sector
5      int flag1 = 1;
6      sector = 0;
7      while(sector < num){
8          char buffer_bitmap[SECTOR_SIZE]; // bitmap size (char because each char is equivalent to 1 byte)
9          memset(buffer_bitmap, 0, SECTOR_SIZE); // fill block of memory with 0
10         for (byte = 0; byte < SECTOR_SIZE; byte++){ // looping each byte
11             if(!flag1) {
12                 break;
13             }
14             else {
15                 for (bit = 7; bit >= 0; bit--){ // looping each bit of every byte
16                     if (nbits-- > 0){ // chekcing nbits after decreament by 1
17                         buffer_bitmap[byte] |= (1 << bit);
18                     }
19                     else{
20                         flag1 = 0;
21                         break;
22                     } // set 1 for first nbits (bitwise OR) for leading the start
23                 }
24             }
25         }
26         //update on the disk
27         Disk_Write(start+sector, buffer_bitmap);
28         sector++;
29     }
30 }

```

Figure 1: Implementation of bitmap initialization

2.1.2 Finding First Unused Bit

To allocate memory space, we need to find the index of first unused bit in the bitmap. The implementation of this function is illustrated in Fig. 2. After pulling the current bitmap from the disk, it is checking for first 0 bit through each bit of every byte. Each bit of the current bitmap is shifted to the right and performing a bitwise AND with a 1 to check whether the examining bit is unused(0) or not. Since it is an AND operation, if the reference bit is 0, the result will be also 0 that means we found the first unused bit and the index of that bit is returned after flipping that bit from 0 to 1 using left shift and bitwise OR operation (line#15-19). Otherwise return -1 if the bitmap is already full indicates no more zeros. We improve the

efficiency of the code by checking the index(id) in line#22-24. More precisely, if id has already exceed nbits, we are not required to check any further which reduces the execution cost of loop.

```

1 static int bitmap_first_unused(int start, int num, int nbits){
2     int id = 0; //starting index is zero by default
3     char check_bit_empty;
4     char buffer_bitmap[SECTOR_SIZE];
5     int bit, byte, sector;
6
7     // checking every sectors to find the first unused memory
8     sector = 0;
9     while(sector < num){
10         Disk_Read(start+sector, buffer_bitmap); // pull bitmap of current content
11         for (byte = 0; byte < SECTOR_SIZE; byte++){ // looping each byte
12             for (bit = 7; bit >= 0; bit--){ // looping every bit
13                 // *****
14                 // Important part
15                 check_bit_empty = (buffer_bitmap[byte] >> bit) & 1; // Shift and ANDing for knowing which bits are still high
16                 if (check_bit_empty == 0){ // if used bit found, set it to 0
17                     buffer_bitmap[byte] |= 1 << bit;
18                     Disk_Write(start+sector, buffer_bitmap); // update the disk
19                     return id; // return index
20                 }
21                 id++; // increment index by 1
22                 if (id > nbits){ // return -1 if the index exceed the maximum number of bitmaps
23                     return -1;
24                 }
25                 // *****
26             }
27         }
28         sector++;
29     }
30     return -1; // return -1 if first unused not found
31 }

```

Figure 2: Implementation of *bitmap_first_unused()*

2.1.3 Reset

When a file/inode is removed, we need to manage released free space. *bitmap_reset()* function is finding the i-th bit location and set it to 0. A single high(1) bit is left shifted then bitwise NOT is applied to perform an AND operation that reset the ibit. See the implementation snapshot in Fig 3 line#11-16.

```

1 static int bitmap_reset(int start, int num, int ibit) {
2     char buffer_bitmap[SECTOR_SIZE];
3     int bit, byte, sector;
4
5     // loop through each sector in bitmap
6     sector = 0;
7     while(sector < num){
8         Disk_Read(start+sector, buffer_bitmap);
9         for (byte = 0; byte < SECTOR_SIZE; byte++){ // checking every byte
10             for (bit = 7; bit >= 0; bit--){ // checking every bit
11                 if(!ibit){ // once our bit is found set to 0 and write changes
12                     buffer_bitmap[byte] &= ~(1 << bit); // Left shift NOT of 1 then AND to set 0
13                     Disk_Write(start+sector, buffer_bitmap);
14                     return 0;
15                 }
16                 ibit--;
17             }
18         }
19         sector++;
20     }
21     return -1;
22 }

```

Figure 3: Implementation of *bitmap_reset()*

2.1.4 Checking File Name

At the time of creating new file, given file name should be suitable for OS file system. The implementation of the function *illegal_file_name()* is shown in Fig. 4. Legal characters for

a file name include letters (case sensitive), numbers, dots, dashes, and underscores. There are some build-in string function is available those are used for this purpose. *isdigit()* checks given character is digit or not, *isalpha()* checks given character is a alphabet or not. Another part of the code from line#14-17 is checking the length of the file. We have used built-in function *strlen()* to check the length of the filename where file name should not be more than MAX_NAME-1. Finally, it returns 1 if file name/length is incorrect otherwise return 0.

```

1  static int illegal_filename(char* name) {
2      int i;
3      char end[] = "-_.";
4
5      //checking character by character
6      for(i = 0; i < strlen(name); i++){
7          if(!isdigit(name[i]) && !isalpha(name[i]) && !strchr(end, name[i])){
8              printf("Illegal character found! %c\n", name[i]);
9              return 1;
10         }
11     }
12
13     //checking length
14     if(strlen(name) >= MAX_NAME-1){
15         printf("Bad Length\n");
16         return 1;
17     }
18
19     //return false otherwise
20     return 0;
21 }

```

Figure 4: Implementation of *illegal_file_name()*

2.2 Read/Write

2.2.1 File Read

File read function (File_Read()) reads specific bytes from a open file which is referenced by a file descriptor (fd). The data is read into a buffer (function given parameter).

Every open file has a pointer which indicated the current location of the file read/write. This function start reading the file from that position and updated the pointer position at the end of reading.

The code checks if the file is open or not. If it is not open then the code returns -1 and set osErrno to E_BAD_FD. If the file is open then the read operation performed.

The function perform the read for a specified size (function parameter) if there is enough data on the file. If the available data is less the the specified size then only the available portion of the data is returned.

The function also check if the file pointer is already at the end of the file. In such case no data is returned. [Figure 5].

2.2.2 File Write

File write function (File_Write()) write specific bytes from buffer to a open file which is referenced by fd.

Every open file has a pointer which indicated the current location of the file read/write. This function start writing the file from that position and updated the pointer position at the end of writing.

The code checks if the file is open or not. If it is not open then the code returns -1 and set osErrno to E_BAD_FD. If the file is open then the write operation performed and value of written file size if return.

```

850 /* YOUR CODE */
851 int File_Read(int fd, void* buffer, int size)
852 {
853     int reading_part = 0; // File is read and kept a copy into reading_part
854
855     open_file_t openFileInfo = open_files[fd]; // Information about the file fd
856
857     // First check the file is open or not
858     if (openFileInfo.inode == 0)
859     {
860         dprintf("...file not open");
861         osErrno = E_BAD_FD;
862         return -1;
863     }
864
865     // load the disk sector containing the inode
866     int inode_sector = INODE_TABLE_START_SECTOR + openFileInfo.inode / INODES_PER_SECTOR;
867     char inode_buffer[SECTOR_SIZE];
868     if (Disk_Read(inode_sector, inode_buffer) < 0) {
869         osErrno = E_GENERAL;
870         return -1;
871     }
872     dprintf("... load inode table for inode from disk sector %d\n", inode_sector);
873
874     // get the inode
875     int inode_start_entry = (inode_sector - INODE_TABLE_START_SECTOR) * INODES_PER_SECTOR;
876     int offset = openFileInfo.inode - inode_start_entry;
877     assert(0 <= offset && offset < INODES_PER_SECTOR);
878     inode_t *child_inode = (inode_t *) (inode_buffer + offset * sizeof(inode_t));
879
880
881     int sector_read = openFileInfo.pos / SECTOR_SIZE; // The sector which need to read
882
883     // Check if the position is at the end of file, or data at that sector does not exists
884     if (openFileInfo.pos == MAX_FILE_SIZE || !child_inode->data[sector_read]) {
885         return 0;
886     }
887
888     char sectorBuffer[SECTOR_SIZE];
889     // check if there any problem to read the data
890     if (Disk_Read(child_inode->data[sector_read], sectorBuffer) < 0) {
891         return -1;
892     }
893     dprintf("... load disk sector %d\n", child_inode->data[sector_read]);
894
895
896     void *reading_pos = sectorBuffer + (open_files[fd].pos - (openFileInfo.pos / SECTOR_SIZE) * SECTOR_SIZE);
897
898     // Read the data from the file until maximum file is already read completely or
899     // requested read size is already read or all the data has been accessed
900     while (open_files[fd].pos < MAX_FILE_SIZE && size > 0 && child_inode->data[sector_read]) {
901
902         if (size >= SECTOR_SIZE) {
903             dprintf("... Reading BYTE %d\n", open_files[fd].pos);
904             memcpy(buffer, reading_pos, SECTOR_SIZE);
905             size -= SECTOR_SIZE;
906             buffer += SECTOR_SIZE;
907             reading_part += SECTOR_SIZE;
908             open_files[fd].pos += SECTOR_SIZE;
909             sector_read++;
910             if (open_files[fd].pos < MAX_FILE_SIZE) { // check if file read reach to maximum file size
911                 if (Disk_Read(child_inode->data[sector_read], sectorBuffer) < 0) {
912                     return -1;
913                 }
914                 reading_pos = sectorBuffer;
915             } else {
916                 return reading_part;
917             }
918         } else {
919             dprintf("... Reading BYTE %d\n", open_files[fd].pos);
920             memcpy(buffer, reading_pos, (size_t) size);
921             reading_part += size;
922             open_files[fd].pos += size;
923             return reading_part;
924         }
925     }
926
927     return reading_part;
928     //return -1;
929 }
930

```

Figure 5: Implementation of *File_Read()*

The code also checks if the write can not be performed because of lack of space or file exceed maximum file size. In case of lack of space, it returns -1 and set osErrno to E_NO_SPACE. if the file exceeds the maximum file size, it also returns -1 and set osErrno to E_FILE_TOO_BIG.[Figure 6].


```

932 /* YOUR CODE */
933 int File_Write(int fd, void* buffer, int size){
934
935     // Have the information about the file (fd) where we have to write the buffer
936     open_file_t openFileInfo = open_files[fd];
937
938     // First check the file is open or not
939     if (openFileInfo.inode == 0) {
940         dprintf("... file not open");
941         osErrno = E_BAD_FD; // if not open set the osError to E_BAD_FD
942         return -1;
943     }
944
945     // check if the file have enough space to write
946     if ((openFileInfo.size + size) > MAX_FILE_SIZE) {
947         dprintf("... file is too big.");
948         osErrno = E_FILE_TOO_BIG;
949         return -1;
950     }
951
952     // load the disk sector containing the inode
953     int inode_sector = INODE_TABLE_START_SECTOR + openFileInfo.inode / INODES_PER_SECTOR;
954     char inode_buffer[SECTOR_SIZE];
955     if (Disk_Read(inode_sector, inode_buffer) < 0) {
956         osErrno = E_GENERAL;
957         return -1;
958     }
959     dprintf("... load inode table for inode from disk sector %d\n", inode_sector);
960
961     // get the inode
962     int inode_start_entry = (inode_sector - INODE_TABLE_START_SECTOR) * INODES_PER_SECTOR;
963     int offset = openFileInfo.inode - inode_start_entry;
964     assert(0 <= offset && offset < INODES_PER_SECTOR);
965     inode_t *fileInode = (inode_t *) (inode_buffer + offset * sizeof(inode_t));
966
967     int sectors_need = (size / SECTOR_SIZE) + 1; // number of sectors needed to write the buffer
968
969     // start the process of writing
970     for (int i = 0; i < sectors_need; i++) {
971         char sector[SECTOR_SIZE];
972         int sector_ID = bitmap_first_unused(SECTOR_BITMAP_START_SECTOR, SECTOR_BITMAP_SECTORS, SECTOR_BITMAP_SIZE);
973
974         // Check if there is any sector left for writing
975         if (sector_ID < 0) {
976             osErrno = E_NO_SPACE;
977             dprintf("... no space left.");
978             Disk_Write(INODE_TABLE_START_SECTOR + (openFileInfo.inode / INODES_PER_SECTOR), inode_buffer);
979             return -1;
980         }
981         fileInode->data[i] = sector_ID;
982         Disk_Read(sector_ID, sector);
983
984         // write buffer to sector. Buffer might be bigger than sector
985         // need to request more sectors
986         if (size > SECTOR_SIZE) {
987             size -= SECTOR_SIZE;
988             memcpy(sector, buffer, SECTOR_SIZE);
989             buffer += SECTOR_SIZE;
990             fileInode->size += SECTOR_SIZE;
991         } else {
992             memcpy(sector, buffer, (size_t) size);
993             buffer += size;
994             fileInode->size += size;
995             size = 0;
996         }
997         openFileInfo.size = fileInode->size;
998         openFileInfo.pos = fileInode->size;
999         Disk_Write(sector_ID, sector);
1000     }
1001     Disk_Write(INODE_TABLE_START_SECTOR + (openFileInfo.inode / INODES_PER_SECTOR), inode_buffer);
1002     return 0;
1003 }

```

Figure 6: Implementation of *File_Write()*

2.2.3 File Seek

File seek function (*File_Seek()*) updates the current location of the file pointer which is provided as an offset from the beginning of the file.

The function first checks if the file is open or not. If it is not open then the code returns -1 and set *osErrno* to *E_BAD_FD*. If the file is open then the file seek operation is performed. This function also checks if the offset is larger than the file size or negative then it returns -1 and set *osErrno* to *E_SEEK_OUT_OF_BOUNDS*. [Figure 7].

```

1006  /* YOUR CODE */
1007  int File_Seek(int fd, int offset){
1008      open_file_t openFileInfo = open_files[fd];
1009
1010      if (openFileInfo.inode == 0) { // check if the file is open or not
1011          dprintf("... file not open");
1012          osErrno = E_BAD_FD;
1013          return -1;
1014      }
1015
1016      //Check if offset is larger than the size of the file or offset is negative
1017      if (offset > openFileInfo.size || offset < 0) {
1018          dprintf("... offset is out of bound");
1019          osErrno = E_SEEK_OUT_OF_BOUNDS;
1020          return -1;
1021      }
1022
1023      open_files[fd].pos = offset;
1024      return open_files[fd].pos;
1025  }

```

Figure 7: Implementation of *File_Seek()*

2.2.4 Directory Read

Directory read function (*Dir_Read()*) reads the contents of a directory and returns the set of directory entries into the buffer.

The function check if the buffer size is large enough to hold all the entries of the directory. If not then the function returns -1 and set *osErrno* to *E_BUFFER_TOO_SMALL*.

If the function successfully read the data into the buffer then it returns the number of directory entries that are in the directory.[Figure 8].

2.2.5 Directory Size

Directory size function (*Dir_Size()*) returns the number of bytes in the directory referred to by path.

This function first check if the directory exists or not. If exists then it also check if it is a directory or file. If it is a directory then the size of the directory is returned. [Figure 9].

2.3 Remove

In order to remove the file/directory, we need to follow similar steps as create file/directory. In particular, we need to remove corresponding iNode, update the iNodes Bitmap, remove the content from the disk, and update the Sector Bitmap.

The following functions are responsible for the object removing:

- *File_Unlink(char* file)* - to remove the file with path *file*
- *Dir_Unlink(char* path)* - to remove the directory with path *path*

Both functions *File_Unlink(char* file)* and *Dir_Unlink(char* path)* (Figure 10) use the helper function *remove_file_or_directory(int type, char* pathname)*, which has similar structure with *create_file_or_directory(int type, char* pathname)*. Variable *type* corresponds to the type of object (0 for files, and 1 for directory). Attribute *pathname* is the absolute path of the file/directory that we want to remove.

remove_file_or_directory(int type, char* pathname)

On Figure 11 you can observe the structure of the function *remove_file_or_directory(int type, char* pathname)*.

```

1100 /* YOUR CODE */
1101 int Dir_Read(char* path, void* buffer, int size) {
1102     int parent_inode; //inode of the directory that need to read
1103     follow_path(path, &parent_inode, NULL);
1104
1105     if (parent_inode >= 0) { //check the directory is found or not
1106         // load the disk sector containing the inode
1107         int inode_sector = INODE_TABLE_START_SECTOR+parent_inode/INODES_PER_SECTOR;
1108         char inode_buffer[SECTOR_SIZE];
1109         if(Disk_Read(inode_sector, inode_buffer) < 0) {
1110             osErrno = E_GENERAL;
1111             return -1;
1112         }
1113         dprintf("... load inode table for inode from disk sector %d\n", inode_sector);
1114
1115         // get the inode
1116         int inode_start_entry = (inode_sector-INODE_TABLE_START_SECTOR)*INODES_PER_SECTOR;
1117         int offset = parent_inode-inode_start_entry;
1118         assert(0 <= offset && offset < INODES_PER_SECTOR);
1119         inode_t* directory_inode = (inode_t*)(inode_buffer+offset*sizeof(inode_t));
1120         dprintf("... inode %d (size=%d, type=%d)\n", parent_inode, directory_inode->size, directory_inode->type);
1121
1122         // check weather the inode type is directory
1123         if(directory_inode->type != 1) {
1124             dprintf("... error: '%s' is not a Directory\n", path);
1125             osErrno = E_GENERAL;
1126             return -1;
1127         }
1128
1129         //check if the read buffer is larger enough to hold the details
1130         if(directory_inode->size*sizeof(dirent_t) > size){
1131             dprintf("ERROR: type %d inode %d size %d givensize %d\n", directory_inode->type, parent_inode, directory_inode->size, size);
1132             osErrno = E_BUFFER_TOO_SMALL;
1133             return -1;
1134         }
1135     }
1136
1137     int sector, i;
1138     int increase_size = 0; // increase the buffer size by this amount
1139     char dirent_buffer[SECTOR_SIZE];
1140
1141     for(sector = 0; sector < MAX_SECTORS_PER_FILE; sector++){ //iterate over each sector
1142         if(directory_inode->data[sector]){
1143             Disk_Read(directory_inode->data[sector], dirent_buffer);
1144             for(i = 0; i < DIRENTS_PER_SECTOR; i++){ // for each directory read it and copy it to the buffer
1145                 dirent_t* dirent = (dirent_t*)(dirent_buffer+i*sizeof(dirent_t));
1146                 if(dirent->inode){
1147                     memcpy(buffer+increase_size, (void*)dirent, sizeof(dirent_t));
1148                     increase_size += sizeof(dirent_t);
1149                 }
1150             }
1151         }
1152     }
1153
1154     dprintf("%d\n", directory_inode->size);
1155     return directory_inode->size;
1156 } else {
1157     dprintf("... directory '%s' is not found\n", path);
1158     return 0;
1159 }
1160 }
1161 }
1162 }
1163 }

```

Figure 8: Implementation of *Dir_Read()*

On line 9 we are finding iNode numbers of the file and it's parent iNode using function *follow_path(char* path, int* last_inode, char* last_fname)*. This function divides the path into tokens by splitting it with "/" character. Then starting from the root iNode, by iterating through each iNode child and comparing the corresponding filename entry in the *dirent* table with the next level in the path, we are finding the iNode number that corresponds to the last file in the path.

If the iNode for a file (lines 14-21) or parent (lines 48-53) can't be found, we are returning "No such file/directory" error.

On lines 26-28, we are checking if the file is currently open. If so, we can't remove it, since other application might overwrite our changes.

If a child and parent iNode is found, and the file is not currently open, we are attempting to remove the iNode (line 32), by calling the function *remove_inode(int type, int parent_inode, int child_inode)*.


```

1061 /* YOUR CODE */
1062 int Dir_Size(char* path){
1063     int parent_inode; //inode of the directory that need to read
1064     follow_path(path, &parent_inode, NULL);
1065
1066     if (parent_inode >= 0) { //check the directory is found or not
1067
1068         // load the disk sector containing the inode
1069         int inode_sector = INODE_TABLE_START_SECTOR+parent_inode/INODES_PER_SECTOR;
1070         char inode_buffer[SECTOR_SIZE];
1071         if(Disk_Read(inode_sector, inode_buffer) < 0) {
1072             osErrno = E_GENERAL;
1073             return -1;
1074         }
1075         dprintf("... load inode table for inode from disk sector %d\n", inode_sector);
1076
1077         // get the inode
1078         int inode_start_entry = (inode_sector - INODE_TABLE_START_SECTOR) * INODES_PER_SECTOR;
1079         int offset = parent_inode - inode_start_entry;
1080         assert(0 <= offset && offset < INODES_PER_SECTOR);
1081         inode_t *directory_inode = (inode_t *) (inode_buffer + offset * sizeof(inode_t));
1082         dprintf("... inode %d (size=%d, type=%d)\n", parent_inode, directory_inode->size, directory_inode->type);
1083
1084         // check weather the inode type is directory
1085         if (directory_inode->type != 1) {
1086             dprintf("... error: '%s' is not a directory\n", path);
1087             osErrno = E_GENERAL;
1088             return -1;
1089         }
1090         dprintf("... RETURNING SIZE: '%d' \n", (int) (directory_inode->size * sizeof(dirent_t)));
1091         return (int) (directory_inode->size * sizeof(dirent_t));
1092     } else {
1093         dprintf("... directory '%s' is not found\n", path);
1094         return 0;
1095     }
1096 }
1097
1098 }

```

Figure 9: Implementation of *Dir_Size()*

```

/* YOUR CODE */
int File_Unlink(char* file)
{
    dprintf("File_Unlink('%s'):\n", file);
    // 0 - code for file type
    return remove_file_or_directory(0, file);
}

/* YOUR CODE */
int Dir_Unlink(char* path)
{
    dprintf("Dir_Unlink('%s'):\n", path);
    // 1 - code for directory type
    return remove_file_or_directory(1, path);
}

```

Figure 10: General structure of functions to remove File or Directory

remove_inode(int type, int parent_inode, int child_inode)

This function takes care of removing the child iNode, removing the link from the parent iNode to child iNode, removing the content of the file, removing iNode from the disk, and updating sector bitmap and iNode bitmap. The screenshot of this function is at Figure 12.

First, we are loading the child iNode from the disk (lines 9-22). Since on previous step we obtained the child iNode number, we know its exact location on a disk. Therefore, we just load the iNode content to the buffer and transform it into iNode type.

Second, we check for errors (lines 28-38): if the type in the argument doesn't match the iNode type or if we are trying to remove the directory that is not empty.

Third, if we are removing the file, we need to remove all its content from the disk (lines 43-60). Since the child iNode contains the addresses of each data block of the file to the disk sector where it belongs to, we just scan through each such sector and set all the bits there to zero. Along with updating the disk, we also set the bitmap entry that corresponds to each sector of the file to zero.

Forth, we are removing the child iNode from the disk (lines 65-67). Those, we are setting all

```

1 int remove_file_or_directory(int type, char* pathname)
2 {
3     //type=0 -file
4     //type=1 - directory
5     // This function is modification of create_file_or_directory() that removes file or directory
6
7     int child_inode;
8     char last_fname[MAX_NAME];
9     int parent_inode = follow_path(pathname, &child_inode, last_fname);
10    if(parent_inode >= 0) {
11        //-----
12        // return error if there are no such files
13        //-----
14        if(child_inode < 0) {
15            dprintf("... file/directory '%s' doesn't exist \n", pathname);
16            if(type==0) {
17                osErrno = E_NO_SUCH_FILE;
18            } else {
19                osErrno = E_NO_SUCH_DIR;
20            }
21            return -1;
22        } else {
23            //-----
24            // Return error if file is currently open
25            //-----
26            if(type==0 && is_file_open(child_inode)){
27                osErrno = E_FILE_IN_USE;
28                return -1;
29            }
30            // Remove iNode
31            //-----
32            } else if(remove_inode(type, parent_inode, child_inode) >= 0){
33                dprintf("... successfully removed file/directory: '%s'\n", pathname);
34                return 0;
35            }
36            // Return error if child inode couldn't be removed
37            //-----
38            } else {
39                dprintf("... error: something wrong with removing file/directory\n");
40                osErrno = E_NO_SUCH_FILE;
41                return -1;
42            }
43        }
44    } else {
45        //-----
46        // Return error if can't find parent iNode
47        //-----
48        if(type==0){
49            osErrno = E_NO_SUCH_FILE;
50        } else {
51            osErrno = E_NO_SUCH_DIR;
52        }
53        return -1;
54    }
55 }

```

Figure 11: Function to remove file or directory

bits of the iNode's sector to zero. Since we freed the disk space, we need to update the iNode bitmap entry that corresponds to newly free sector (line 72)

Finally, we need to remove the pointer from the parent iNode to the child iNode. For that, we load the parent iNode from the disk (lines 78-89) in the similar maner that we did for the child iNode. Then, we iterate through each data pointer of the parent iNode, and check if it points to the child iNode. If it is, than we remove this pointer from the parent iNode, and we also remove the corresponding entry in a dirent table that maps the child iNode number with it's file name. (lines 100-125).

3 Testing

We have created some custom test cases to show the basic functionalities of the implemented file system. The rusult of make command is shown in Fig. 13. Before running custom test cases, we run simple-test.exe to make sure that our code is working fine. Results of simple test is shown in Fig. 14 that ensures that the following functions are working fine.

- bitmap initialization, finding first unused bit, and reset
- illegal file name
- Creating file/directory

```

1 int remove_inode(int type, int parent_inode, int child_inode)
2 {
3     // This function is a modification to old_inode() that removes child inode
4     dprintf("... removing inode %d\n", child_inode);
5     // Load child inode from the disk
6     // Load the disk sector containing the child inode
7     int inode_sector = INODE_TABLE_START_SECTOR+child_inode/INOES_PER_SECTOR;
8     char inode_buffer[SECTOR_SIZE];
9     if(Disk_Read(inode_sector, inode_buffer) < 0){
10         dprintf("Error: can't read inode from disk sector");
11         return -1;
12     }
13     //dprintf("%s\n", inode_buffer);
14     dprintf("... load inode table for child inode from disk sector %d\n", inode_sector);
15     // get the child inode
16     int inode_start_entry = (inode_sector-INODE_TABLE_START_SECTOR)*INOES_PER_SECTOR;
17     int offset = child_inode-inode_start_entry;
18     assert(0 <= offset && offset < INOES_PER_SECTOR);
19     inode_t* child = (inode_t*)(inode_buffer+offset*sizeof(inode_t));
20     // Check for errors
21     // Check if the type is right
22     if(child->type != type){
23         dprintf("Error: the actual type of inode doesn't match with parameter specified in function remove_inode()");
24         return -3;
25     }
26     //Check if directory is not empty
27     if(child->type == 1 && child->size == 0){
28         dprintf("Error: directory is not empty");
29         osErrno = E_DIR_NOT_EMPTY;
30         return -2;
31     }
32     // Remove the file content from the disk
33     // remove data from disk sectors
34     char buf_sector[SECTOR_SIZE];
35     int sector_address;
36     if(child->type == 1){
37         if(child->size == 0){
38             dprintf("... Removing file data from the disk");
39             for(int i = 0; i < MAX_SECTORS_PER_FILE; i++){
40                 sector_address = child->data[i];
41                 if(sector_address > 0){
42                     if(Disk_Read(sector_address, buf_sector) < 0) return -1; //try to read the sector
43                     memset(buf_sector, 0, SECTOR_SIZE); //set buffer to 0
44                     Disk_Write(sector_address, buf_sector); //write 0s to the sector
45                     bitmap_reset(SECTOR_BITMAP_START_SECTOR, SECTOR_BITMAP_SECTORS, sector_address); //reset bitmap for the current sector
46                 }
47             }
48         } else{
49             dprintf("... The directory/file is empty");
50         }
51     }
52 }

```

(a) Loading child iNode and removing the content from the disk

```

59 // Remove inode from the disk sector
60 //
61 dprintf("... Removing inode from the disk");
62 memset(child, 0, sizeof(inode_t)); //Delete child inodes from memory
63 //memset(inode_buffer, 0, SECTOR_SIZE);
64 if(Disk_Write(inode_sector, inode_buffer) < 0) return -1;
65 // Reset inode bitmap
66 //
67 bitmap_reset(INODE_BITMAP_START_SECTOR, INODE_BITMAP_SECTORS, child_inode); //reset i-node bitmap
68 // Remove link from parent inode to child inode
69 // Read the parent inode from the disk
70 //
71 // get the disk sector containing the parent inode
72 inode_sector = INODE_TABLE_START_SECTOR+parent_inode/INOES_PER_SECTOR;
73 if(Disk_Read(inode_sector, inode_buffer) < 0) return -1;
74 dprintf("... load inode table for parent inode %d from disk sector %d\n",
75     parent_inode, inode_sector);
76 // get the parent inode
77 inode_start_entry = (inode_sector-INODE_TABLE_START_SECTOR)*INOES_PER_SECTOR;
78 offset = parent_inode-inode_start_entry;
79 assert(0 <= offset && offset < INOES_PER_SECTOR);
80 inode_t* parent = (inode_t*)(inode_buffer+offset*sizeof(inode_t));
81 dprintf("... get parent inode %d (size=%d, type=%d)\n",
82     parent_inode, parent->size, parent->type);
83 // Check if parent inode is a directory
84 //
85 if(parent->type != 1){
86     dprintf("... error: parent inode is not directory\n");
87     return -2;
88 }
89 // Update parent inode and dirent table
90 //
91 int nentries = parent->size; // remaining number of directory entries
92 int idx = 0;
93 int found_flag = 0;
94 while(nentries > 0){
95     char dirent_buffer[SECTOR_SIZE]; // cached content of directory entries
96     if(Disk_Read(parent->data[idx], dirent_buffer) < 0) return -2;
97     for(int i = 0; i < DIRENTS_PER_SECTOR; i++){
98         if(i < nentries) break;
99         dirent_t* tmp_dirent = (dirent_t*)(dirent_buffer+i*sizeof(dirent_t));
100         if(tmp_dirent->inode == child_inode){ //dirent found
101             dprintf("... Corresponding dirent found\n");
102             char dirent_buffer2[SECTOR_SIZE];
103             int group = parent->size/DIRENTS_PER_SECTOR;
104             if(Disk_Read(parent->data[group], dirent_buffer2) < 0) return -2;
105             int start_entry = group*DIRENTS_PER_SECTOR;
106             offset = parent->size-start_entry;
107             dirent = parent->data[start_entry];
108             tmp_dirent1 = (dirent_t*)(dirent_buffer2+offset*sizeof(dirent_t));
109             strncpy(tmp_dirent->fname, tmp_dirent1->fname, MAX_NAME);
110             tmp_dirent->inode = tmp_dirent1->inode;
111             memset(tmp_dirent, 0, sizeof(dirent_t));
112             parent->size--;
113             if(Disk_Write(inode_sector, inode_buffer) < 0) return -1; //update parent inode
114             if(Disk_Write(parent->data[group], dirent_buffer2) < 0) return -1; //update dirent
115             if(Disk_Write(parent->data[idx], dirent_buffer)) return -1; //update dirent
116             found_flag = 1;
117         }
118     }
119     idx++; nentries -= DIRENTS_PER_SECTOR;
120 }
121 if(found_flag == 1) return 0;
122 else{
123     dprintf("... could not find dirent\n");
124     return -1; // not found
125 }
126 }

```

(b) Resetting bitmaps, removing child iNode from the disk, updating parent iNode and dirent table

Figure 12: Function to remove iNode

```
(base) mamun@MDS-MacBook-Pro:~/FIU/Fall19/OS/Project2/COP5614-FinalProject$ make
gcc -O -Wall -c main.c -o main.o
gcc -o main.exe main.o -L. -lFS -lDisk
gcc -O -Wall -c simple-test.c -o simple-test.o
gcc -o simple-test.exe simple-test.o -L. -lFS -lDisk
gcc -O -Wall -c slow-ls.c -o slow-ls.o
gcc -o slow-ls.exe slow-ls.o -L. -lFS -lDisk
gcc -O -Wall -c slow-mkdir.c -o slow-mkdir.o
gcc -o slow-mkdir.exe slow-mkdir.o -L. -lFS -lDisk
gcc -O -Wall -c slow-rmdir.c -o slow-rmdir.o
gcc -o slow-rmdir.exe slow-rmdir.o -L. -lFS -lDisk
gcc -O -Wall -c slow-touch.c -o slow-touch.o
gcc -o slow-touch.exe slow-touch.o -L. -lFS -lDisk
gcc -O -Wall -c slow-rm.c -o slow-rm.o
gcc -o slow-rm.exe slow-rm.o -L. -lFS -lDisk
gcc -O -Wall -c slow-cat.c -o slow-cat.o
gcc -o slow-cat.exe slow-cat.o -L. -lFS -lDisk
gcc -O -Wall -c slow-import.c -o slow-import.o
gcc -o slow-import.exe slow-import.o -L. -lFS -lDisk
gcc -O -Wall -c slow-export.c -o slow-export.o
gcc -o slow-export.exe slow-export.o -L. -lFS -lDisk
rm simple-test.o slow-cat.o slow-mkdir.o slow-rm.o slow-rmdir.o slow-import.o slow-touch.o main.o slow-ls.o slow-export.o
```

Figure 13: Output of 'make'

- Removing file/directory/inode

We design a custom test cases as similar as Fig. 15 where three directories and two files will be created under root '/'. Also, dir2 will have two sub directories such as dir21, and dir22, finally dir21 will have two files file21a and file21b. We run the following command

- ./slow-mkdir.exe /dir1
- ./slow-touch.exe /file1
- ./slow-mkdir.exe /dir2
- ./slow-touch.exe /file2
- ./slow-mkdir.exe /dir3
- ./slow-mkdir.exe /dir2/dir21
- ./slow-mkdir.exe /dir2/dir22
- ./slow-touch.exe /dir2/dir21/file21a
- ./slow-touch.exe /dir2/dir21/file21b

After running the above lines of command, directories and files has been created and the current inode structure is showing in Fig 15. Now, we will delete '/file1' and add '/dir0' to see inode. Fig. 16 confirms that file1 has been deleted and fig. 17 confirms that unused inode-2 is now using by dir0.

The slow-cat.exe, slow-export.exe, and slow-import.exe test scenarios are also implemented with our code and found that those are working perfectly. [Figure 18].

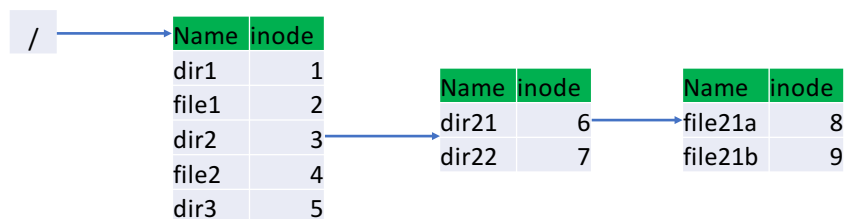
4 Project challenges

The main challenge for this project was to implement `remove_inode()` correctly. We encounter a problem that after removing the file or directory, the newly created object was getting assigned to the end of the dirent table of the parent directory, creating the segmentation. As a result, the `inode->size` does not represent the actual size of the directory list. The remove inode function is modified to overcome the segmentation. Instead of deleting the entries in the middle, the function copied the last entry into the blank position.[Figure 19].

```
(base) mamun@MDs-MacBook-Pro:~/FIU/Fall19/OS/Project2/COP5614-FinalProject$ ./simple-test.exe default-disk
file system booted from file 'default-disk'
file '/first-file' created successfully
file '/second-file' created successfully
dir '/first-dir' created successfully
dir '/first-dir/second-dir' created successfully
ERROR: can't create dir '/first-file/second-dir'
Illegal character found! *
ERROR: can't create dir '/first_dir/third*dir'
file '/first-file' unlinked successfully
ERROR: can't unlink dir '/first-dir'
dir '/first-dir/second-dir' unlinked successfully
file '/second-file' opened successfully, fd=0
successfully wrote 1024 bytes to fd=0
fd 0 closed successfully
file system sync'd to file 'default-disk'
```

Figure 14: Output of simple-test

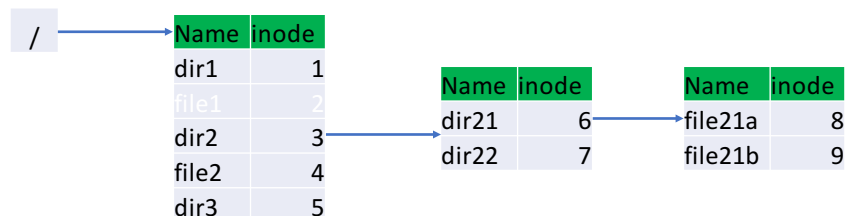
A. Initial inode structure



```
(base) mamun@MDs-MacBook-Pro:~/FIU/Fall19/OS/Project2/COP5614-FinalProject$ ./slow-ls.exe /
directory '/':
  NAME      INODE
0  dir1      1
1  file1     2
2  dir2      3
3  file2     4
4  dir3      5
(base) mamun@MDs-MacBook-Pro:~/FIU/Fall19/OS/Project2/COP5614-FinalProject$ ./slow-ls.exe /dir2
directory '/dir2':
  NAME      INODE
0  dir21     6
1  dir22     7
(base) mamun@MDs-MacBook-Pro:~/FIU/Fall19/OS/Project2/COP5614-FinalProject$ ./slow-ls.exe /dir2/dir21
directory '/dir2/dir21':
  NAME      INODE
0  file21a   8
1  file21b   9
```

Figure 15: Initial inode structure

B. After removing '/file1'



```
(base) mamun@MDs-MacBook-Pro:~/FIU/Fall19/OS/Project2/COP5614-FinalProject$ ./slow-rm.exe /file1
file '/file1' removed successfully
(base) mamun@MDs-MacBook-Pro:~/FIU/Fall19/OS/Project2/COP5614-FinalProject$ ./slow-ls.exe /
directory '/':
  NAME      INODE
0  dir1      1
1  dir2      3
2  file2     4
3  dir3      5
```

Figure 16: inode structure after removing '/file1'

C. After adding '/dir0'

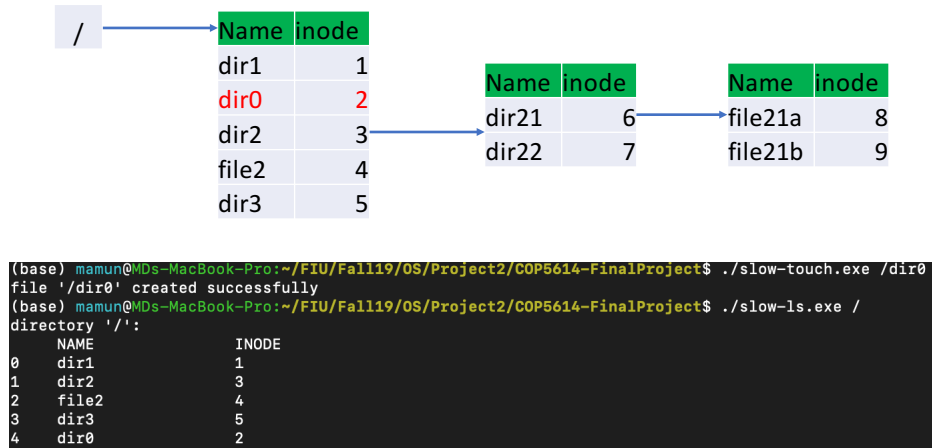


Figure 17: inode structure after adding '/dir0'

```
Shahadats-MacBook-Air:OS_Final_Project shahadat$ ./slow-import.exe /f1.txt SampleText.txt
Shahadats-MacBook-Air:OS_Final_Project shahadat$ ./slow-cat.exe /f1.txt
This is a test file.

This is a test file.

End of File.
```

Figure 18: Output of slow-import and slow-cat

```
333 int nentries = parent->size; // remaining number of directory entries
334 int idx = 0;
335 int found_flag=0;
336 while(nentries > 0) {
337     char dirent_buffer[SECTOR_SIZE]; // cached content of directory entries
338     if(Disk_Read(parent->data[idx], dirent_buffer) < 0) return -2;
339     for(int i=0; i<DIRENTS_PER_SECTOR; i++) {
340         if(i>nentries) break;
341         dirent_t* tmp_dirent = (dirent_t*)(dirent_buffer + i*sizeof(dirent_t));
342
343         if(tmp_dirent->inode==child_inode){//dirent found
344             dprintf("... Corresponding dirent found\n");
345             char dirent_buffer2[SECTOR_SIZE];
346
347             int group = parent->size/DIRENTS_PER_SECTOR;
348             if(Disk_Read(parent->data[group], dirent_buffer2) < 0) return -2;
349             int start_entry = group*DIRENTS_PER_SECTOR;
350             offset = parent->size-start_entry;
351             dirent_t* tmp_dirent1 = (dirent_t*)(dirent_buffer2+offset*sizeof(dirent_t));
352
353             strncpy(tmp_dirent->fname, tmp_dirent1->fname, MAX_NAME);
354             tmp_dirent->inode = tmp_dirent1->inode;
355             memset(tmp_dirent1, 0, sizeof(dirent_t));
356             parent->size--;
357             if(Disk_Write(inode_sector, inode_buffer)<0) return -1;//update parent inode
358             if(Disk_Write(parent->data[group], dirent_buffer2)) return -1;//update dirent
359             if(Disk_Write(parent->data[idx], dirent_buffer)) return -1;//update dirent
360             found_flag=1;
361         }
362     }
363     idx++; nentries -= DIRENTS_PER_SECTOR;
364 }
365 if(found_flag==1) return 0;
366 else{
367     dprintf("... could not find dirent\n");
368     return -1; // not found
369 }
```

Figure 19: Modification in *Remove_Inode()* function

5 Conclusion

In this project, we have developed a file management system that can perform some actions such as read or write of a file or directory, directory list, create a new file or directory, and delete a file or directory. The whole file system is managed by a bitmap system which is also designed here. The test cases shows that the functions works perfectly.